

NGT resource allocation

The NGT GPU allocation problem, measured 30 days of cluster telemetry, and what a policy would change

Felice Pantaleo (CERN)
July 2026

Method

Data

Pod lifecycle, requests, placement, cordons and per-GPU utilization from cluster monitoring, 30 days at 5 min resolution.

Scope

1571 GPU/MIG requests from 114 users on the on-premise pools. STEAM Academy accounts and cloud T4 nodes excluded.

Attribution

Users mapped to WP1 to WP4 via the project roster, department rules and manual classification: 99.4% of GPU-hours attributed.


A request = one pod asking for GPUs. Wait = creation to scheduling.

Idle = GPU utilization below 5%.

NGT resource allocation

Static allocation


- All pools at their effective ceiling around the clock
- No diurnal variation: allocation follows ownership, not workload
- Only the forced maintenance evictions of 6 to 8 July freed capacity



NGT cluster, 18 Jun to 18 Jul 2026. Felice Pantaleo (CERN)
NVL: MIG-enabled GPUs are not visible to DCGM, so the full-GPU ceiling sits below the 96-GPU capacity line. AMD from user-pod accounting (no DCGM).

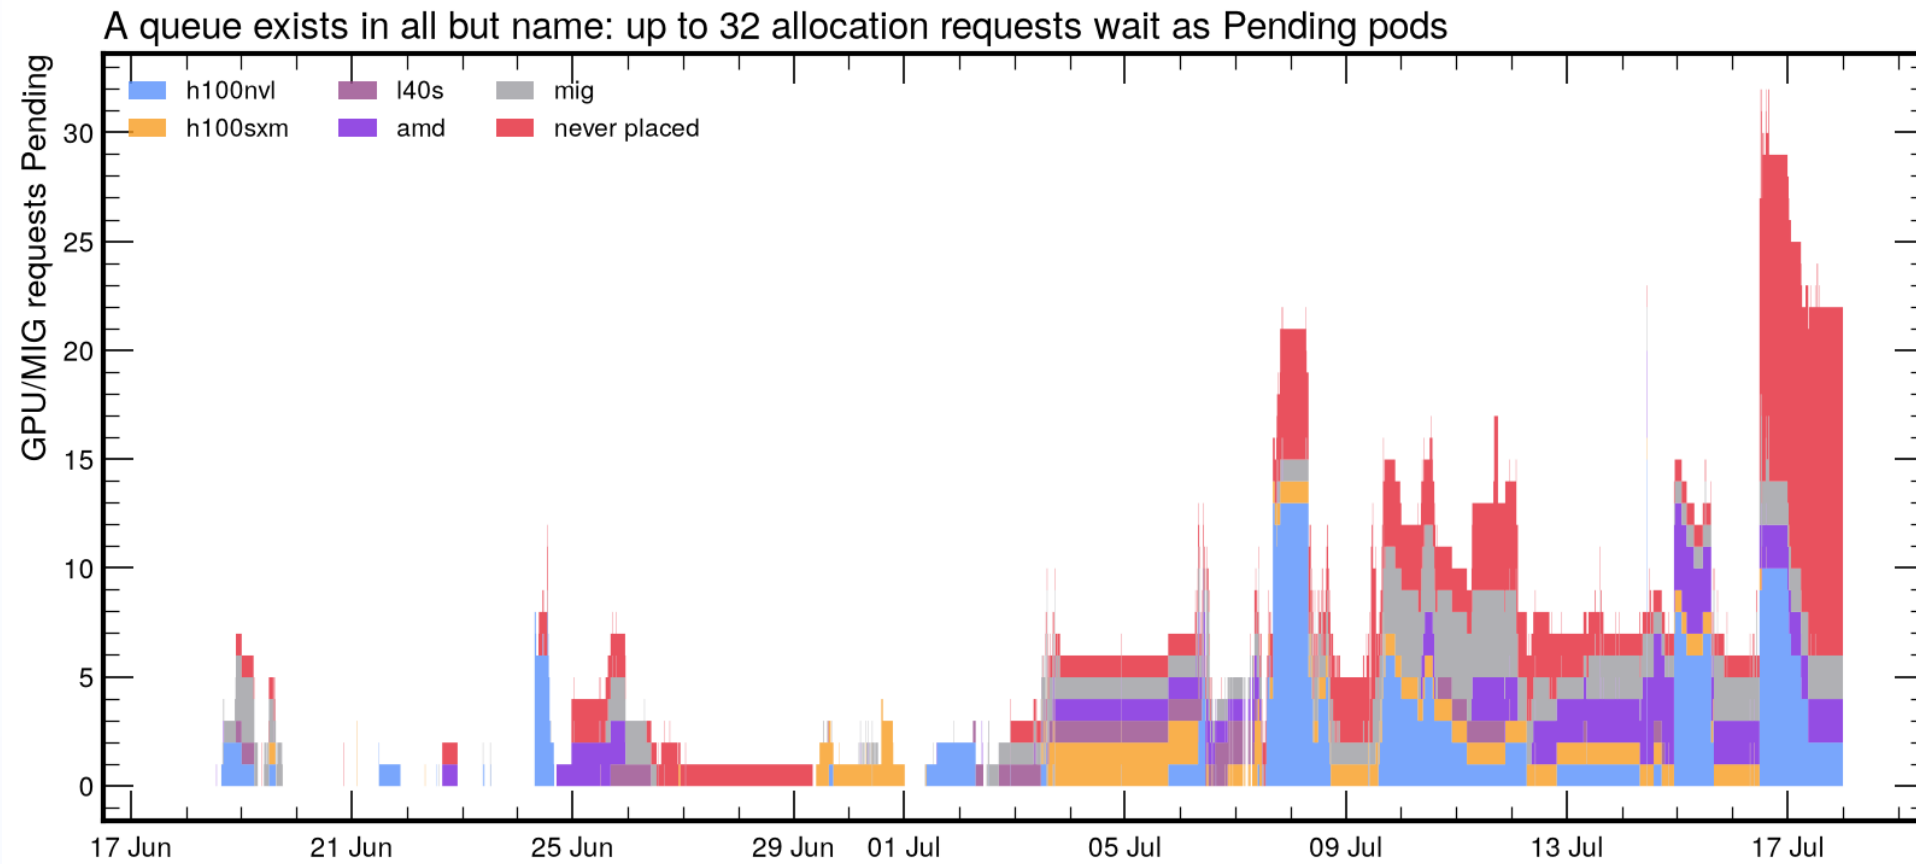
NGT resource allocation

Waiting times




NGT resource allocation

The Pending queue



NGT cluster, 18 Jun to 18 Jul 2026. Felice Pantaleo (CERN)

Lockouts




Over half of all queue pressure is top-up demand from users already running.

NGT resource allocation


Idle held GPU-hours

- 53 000 GPU-hours parked in 30 days
- 267 kCHF/month cloud equivalent
- Right axes: kCHF/day per pool
- Continuous idle baseline on NVL and L40S; bursty parked episodes on SXM



DCGM utilization available for 500 full-GPU allocations; MIG slices and AMD not covered. Cloud rates as in the cost plot. NGT cluster, 18 Jun to 18 Jul 2026. Felice Pantaleo (CERN)


Idle share by pool



NGT resource allocation

Heavy idle holders


- Solid = held idle, pale = active
- Dashed line: one dev GPU held 24/7 all month (720 GPU-h), the legitimate P1 allowance
- **20 users** exceed it
- **26 700 GPU-hours** sit above it
- The consumption and idleness rankings select different users: the top consumer is almost entirely active



NGT resource allocation


Idle holders by pool

- NVL: broad behavior, top 12 hold 66% of idle
- SXM: three users own the idle hours
- L40S: 12 users own 98% of the idle hours
- The one-session rule and batch-only multi-GPU bound all of this by construction




NGT resource allocation

Cloud-equivalent cost



NGT resource allocation

GPU-hours by WP



NGT PMC, July 2026
WP2 is on its 30% target organically; WP3 runs at 47% (CMS production

The real month, replayed

Same 30 days, same requests, replayed through the proposed policy:

	today (FCFS)	proposed
wait p95	557 min	0 min
requests satisfied	96.7%	99.4%
dev sessions within 15 min	91%	98%
NVL idle-held GPU-hours	35 800	13 600
running work terminated by the system	none	none

Policy: one interactive session per member (a new session supersedes the previous one); multi-GPU beyond a 96 GPU-h/month interactive allowance runs as batch behind a WP fair-share priority queue.
Fixed-behavior replay, identical requests for every policy; the FCFS baseline reproduces observed waits (median exact, p95 within factor 2).

Proposal to the PMC

One interactive session

One single-GPU session per member, guaranteed and always available; opening a new one replaces the old.

Multi-GPU is batch

Submitted, executed, exits at completion, behind a priority queue with WP fair share. A 96 GPU-h/month interactive allowance covers debugging.


Account and adjust

GPU time charged to WPs with model factors. All parameters simulated and tunable on the real trace.


NGT resource allocation

Backup

Backup: unsatisfied requests




Backup: hold durations



NGT resource allocation

Backup: top consumers per pool

- Secondary axis: cloud-equivalent kCHF, cheapest of AWS/GCP, July 2026
- Dashed: one GPU 24/7 all month
- 31 users held more than one GPU-month



NGT resource allocation

Backup: pools and cordons

